

DTU



Perception for Autonomous Systems 31392:

# Visual SLAM

## *Simultaneous Localization & Mapping*


Lecturer: Evangelos Boukas—PhD

# Visual Odometry is great

- Lets Sum Up What we did last time
- We performed 3D-to-2D

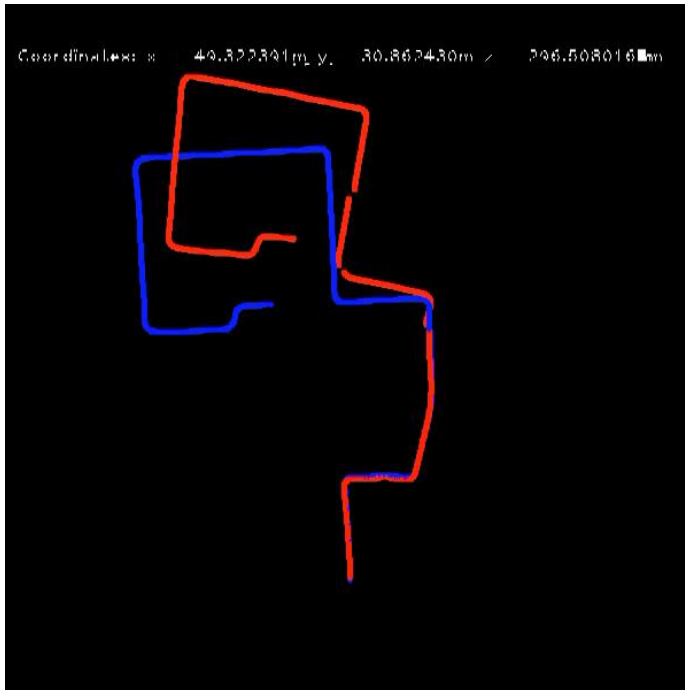
## Algorithm 3. VO from 3-D-to-2-D Correspondences.

- 1) Do only once:
  - 1.1) Capture two frames  $I_{k-2}, I_{k-1}$
  - 1.2) Extract and match features between them
  - 1.3) Triangulate features from  $I_{k-2}, I_{k-1}$
- 2) Do at each iteration:
  - 2.1) Capture new frame  $I_k$
  - 2.2) Extract features and match with previous frame  $I_{k-1}$
  - 2.3) Compute camera pose (PnP) from 3-D-to-2-D matches
  - 2.4) Triangulate all new feature matches between  $I_k$  and  $I_{k-1}$
  - 2.5) Iterate from 2.1).



# Visual Odometry is great

- Lets Sum Up What we did last time
- We performed 3D-to-2D





# Visual Odometry is great

- Lets Sum Up What we did last time
- We performed 3D-to-2D
- Any better approach? How could we



# Visual Odometry is great

- Lets Sum Up What we did last time
- We performed 3D-to-2D
- Any better approach? How could we
  - How about some more features?

# Visual Odometry is great

- Lets Sum Up What we did last time
- We performed 3D-to-2D
- Any better approach? How could we
  - How about some more features?
    - Let's keep the features from multiple frames

# Visual Odometry is great

- Lets Sum Up What we did last time
- We performed 3D-to-2D
- Any better approach? How could we
  - How about some more features?
    - Let's keep the features from multiple frames
  - How could do we know that those are new features?




# Visual Odometry is great

- Lets Sum Up What we did last time
- We performed 3D-to-2D
- Any better approach? How could we
  - How about some more features?
    - Let's keep the features from multiple frames
  - How could do we know that those are new features?

# Visual Odometry is great


- Lets Sum Up What we did last time
- We performed 3D-to-2D

Coordinates: x: 53.169407m, y: 1.180850m, z: 0.37474980m



# Visual Odometry is great



- Anything more we mentioned?
  - Windowed Bundle Adjustment (BA)




- Similar to pose-optimization but it also optimizes 3D points
- In order to not get stuck in local minima, the initialization should be close the minimum
- Levenberg-Marquadt can be used

- Visual Odometry
- Structure from Motion (SFM)
- Bundle Adjustment
- Visual SLAM


- Visual Odometry
- Structure from Motion (SfM)
- Bundle Adjustment
- Visual SLAM



- Visual Odometry
- Structure from Motion (SFM)
- Bundle Adjustment
- Visual SLAM



- Visual Odometry
- Structure from Motion (SfM)
- Bundle Adjustment
- Visual SLAM

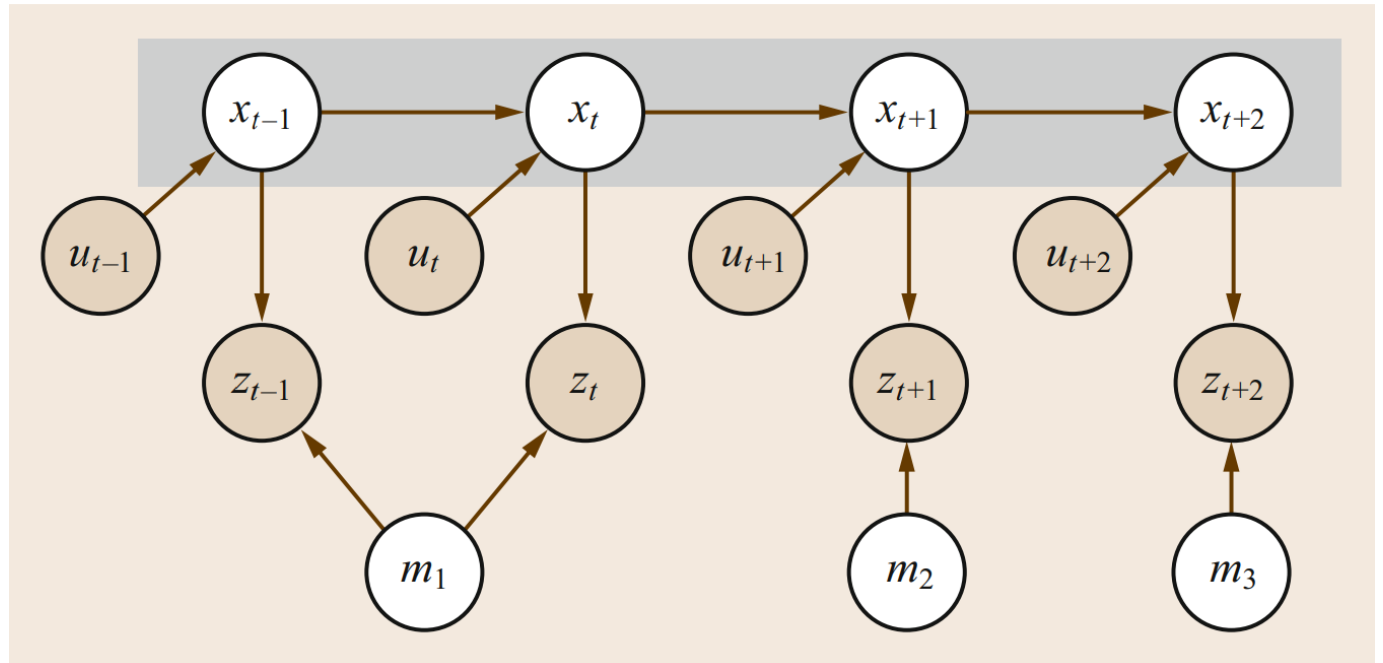


- Sum up Localization from last time
- Some terminology
- Pose-Landmark Graph Slam
- Example of Linear 1D SLAM
- Non-Linear Optimization approaches
- Bundle Adjustment
- Visual Slam System architecture
- ORBSLAM



# Pose-Landmark Graph-Slam

- SLAM problem depicted as Bayes network graph
- At each **location**  $\mathbf{x}_t$
- Observes a nearby feature in the **map**  $\mathbf{m} = \{\mathbf{m}_1; \mathbf{m}_2; \mathbf{m}_3\}$
- Movement  $\mathbf{u}_t$
- An arrow defines causal relationship





# Graph-Based SLAM

## Definition


- Use a graph to represent the problem
- Nodes represent:
  - poses or
  - locations
- Edges Represent:
  - Landmark observations
  - Odometry Measurements
- The minimization optimizes the landmark locations and robot poses

**Graph-Based SLAM: Build the graph and find a node configuration that minimize the error introduced by the constraints**



# Graph-Based SLAM

- Observing previously seen areas generates constraints between non-successive poses
- Treat constraints (motion and measurement) as “soft” elastic springs
- Want to minimize the total energy in the springs



# Graph-Based SLAM

- Observing previously seen areas generates constraints between non-successive poses
- Treat constraints (motion and measurement) as “soft” elastic springs
- Want to minimize the total energy in the springs

We can define the error as follows

- Expected observation (2D sensor):
$$\hat{\mathbf{z}}_{ij}(\underset{\substack{\uparrow \\ \text{robot}}}{\mathbf{x}_i}, \underset{\substack{\uparrow \\ \text{landmark}}}{\mathbf{x}_j}) = \mathbf{R}_i^T(\underset{\substack{\uparrow \\ \text{robot translation}}}{\mathbf{x}_j} - \mathbf{t}_i)$$
  - With the error:  $\mathbf{e}_{ij}(\mathbf{x}_i, \mathbf{x}_j) = \hat{\mathbf{z}}_{ij} - \mathbf{z}_{ij}$ 
$$= \mathbf{R}_i^T(\mathbf{x}_j - \mathbf{t}_i) - \mathbf{z}_{ij}$$

# 1D Linear SLAM

- In the linear case we can solve as follows:

- First construct all constrains

- Absolute Constrains:

$$X(0) = Q - \text{starting position}$$

- Movement Constrains:

$$X(t) = X(t-1) + Dx(t)$$

- Measurement constrains:

$$L(k) = X(t) + N$$

- *Then, solve linear equations*



X0



X1



X2



X3




L0

# 1D Linear SLAM – case 1

- Case 1 - Exact solution exists:

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} -3 \\ 2 \\ 5 \end{bmatrix}$$




# 1D Linear SLAM – case 1

- Case 1 - Exact solution exists:

$$\begin{bmatrix} \phantom{0} & \phantom{0} & \phantom{0} \\ \phantom{0} & \phantom{0} & \phantom{0} \\ \phantom{0} & \phantom{0} & \phantom{0} \end{bmatrix} \cdot \begin{bmatrix} \phantom{0} \\ \phantom{0} \\ \phantom{0} \end{bmatrix} = \begin{bmatrix} \phantom{0} \\ \phantom{0} \\ \phantom{0} \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 0 \\ -1 & 1 & 0 \\ 0 & -1 & 1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \end{bmatrix} \quad A * X = B$$



# 1D Linear SLAM – case 1

- Case 1 - Exact solution exists:


$$\begin{bmatrix} \phantom{0} & \phantom{0} & \phantom{0} \\ \phantom{0} & \phantom{0} & \phantom{0} \\ \phantom{0} & \phantom{0} & \phantom{0} \end{bmatrix} \cdot \begin{bmatrix} \phantom{0} \\ \phantom{0} \\ \phantom{0} \end{bmatrix} = \begin{bmatrix} \phantom{0} \\ \phantom{0} \\ \phantom{0} \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 0 \\ -1 & 1 & 0 \\ 0 & -1 & 1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \end{bmatrix}$$

$$A * X = B$$

$$X = A^{-1} * B$$

$$A^{-1} = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$





# 1D Linear SLAM – case 1

- Case 1 - Exact solution exists:


$$\begin{bmatrix} \phantom{0} & \phantom{0} & \phantom{0} \\ \phantom{0} & \phantom{0} & \phantom{0} \\ \phantom{0} & \phantom{0} & \phantom{0} \end{bmatrix} \cdot \begin{bmatrix} \phantom{0} \\ \phantom{0} \\ \phantom{0} \end{bmatrix} = \begin{bmatrix} \phantom{0} \\ \phantom{0} \\ \phantom{0} \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 0 \\ -1 & 1 & 0 \\ 0 & -1 & 1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \end{bmatrix}$$

$$A * X = B \quad X = A^{-1} * B$$


$$A^{-1} = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

$$x = \begin{bmatrix} -3 & 2 & 5 \end{bmatrix}$$



# 1D Linear SLAM – case 2


- Case 2 – Overdefined problem:
  - $X_0$  sees  $L_0$  at distance 10
  - $X_1$  sees  $L_0$  at distance 5
  - $X_2$  sees  $L_0$  at distance 2



$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} -3 \\ 2 \\ 5 \end{bmatrix}$$

# 1D Linear SLAM – case 2

- Case 2 – Overdefined problem:
  - $X_0$  sees  $L_0$  at distance 10
  - $X_1$  sees  $L_0$  at distance 5
  - $X_2$  sees  $L_0$  at distance 2




$$\begin{bmatrix} \cdot \\ \cdot \\ \cdot \end{bmatrix} \cdot \begin{bmatrix} \cdot \\ \cdot \\ \cdot \end{bmatrix} = \begin{bmatrix} \cdot \\ \cdot \\ \cdot \end{bmatrix}$$

$$A * X = B \quad X = A^{-1} * B \quad A^{-1} = [?]$$

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & -1 & 0 & 0 \\ 0 & 1 & -1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L_0 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \\ -10 \\ -5 \\ -2 \end{bmatrix}$$

# 1D Linear SLAM – case 2

- Case 2 – Overdefined problem:
  - $X_0$  sees  $L_0$  at distance 10
  - $X_1$  sees  $L_0$  at distance 5
  - $X_2$  sees  $L_0$  at distance 2



$$\begin{bmatrix} \phantom{0} & \phantom{0} & \phantom{0} & \phantom{0} \end{bmatrix} \cdot \begin{bmatrix} \phantom{0} & \phantom{0} & \phantom{0} & \phantom{0} \end{bmatrix} = \begin{bmatrix} \phantom{0} & \phantom{0} & \phantom{0} & \phantom{0} \end{bmatrix}$$


$$A * X = B \quad X = A^{-1} * B \quad A^{-1} = [?]$$

$$X = (A^T * A)^{-1} * A^T * B$$

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & -1 & 0 & 0 \\ 0 & 1 & -1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L_0 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \\ -10 \\ -5 \\ -2 \end{bmatrix}$$

# 1D Linear SLAM – case 2

- Case 2 – Overdefined problem:
  - $X_0$  sees  $L_0$  at distance 10
  - $X_1$  sees  $L_0$  at distance 5
  - $X_2$  sees  $L_0$  at distance 2



$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & -1 & 0 & 0 \\ 0 & 1 & -1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L_0 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \\ -10 \\ -5 \\ -2 \end{bmatrix}$$

$$A * X = B \quad X = A^{-1} * B \quad A^{-1} = [?]$$

$$X = (A^T * A)^{-1} * A^T * B$$


$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & -1 & 0 & 0 \\ 0 & 1 & -1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L_0 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \\ -10 \\ -5 \\ -2 \end{bmatrix}$$

$$x = \begin{bmatrix} -3 & 2 & 5 & 7 \end{bmatrix}$$

We infer a consistent landmark position


# 1D Linear SLAM – case 3

- Case 3 – Inconsistent Measurements :
  - $X_0$  sees  $L_0$  at distance 10
  - $X_1$  sees  $L_0$  at distance 5
  - $X_2$  sees  $L_0$  at distance 1 (**Wrong**)



# 1D Linear SLAM – case 3

- Case 3 – Inconsistent Measurements :
  - $X_0$  sees  $L_0$  at distance 10
  - $X_1$  sees  $L_0$  at distance 5
  - $X_2$  sees  $L_0$  at distance 1 (**Wrong**)




$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & -1 & 0 & 0 \\ 0 & 1 & -1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L_0 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \\ -10 \\ -5 \\ -1 \end{bmatrix}$$

$$X = (A^T * A)^{-1} * A^T * B$$

# 1D Linear SLAM – case 3

- Case 3 – Inconsistent Measurements :
  - $X_0$  sees  $L_0$  at distance 10
  - $X_1$  sees  $L_0$  at distance 5
  - $X_2$  sees  $L_0$  at distance 1 (**Wrong**)



$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & -1 & 0 & 0 \\ 0 & 1 & -1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L_0 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \\ -10 \\ -5 \\ -1 \end{bmatrix}$$

$$X = (A^T * A)^{-1} * A^T * B$$

$$x = [-3 \ 2.125 \ 5.5 \ 6.875]$$

We handled inaccurate measurements




# 1D Linear SLAM – case 4

Case 4 – Inconsistent Measurements with Confidence Matrix:

- Linear Least Squares allows us to include a weighting of each linear constraint.
- We can include weights in the computation
- We weight each constraint by a diagonal matrix where the weights are  $1/\text{variance}$  for each constraint.
- Let's say  $X_2$  **variance is 0.2**

$$X = (A^T * W * A)^{-1} * A^T * W * B$$




# 1D Linear SLAM – case 4

Case 4 – Inconsistent Measurements with Confidence Matrix:

- Linear Least Squares allows us to include a weighting of each linear constraint.
- We can include weights in the computation
- We weight each constraint by a diagonal matrix where the weights are  $1/\text{variance}$  for each constraint.
- Let's say  $X_2$  **variance is 0.2**

$$X = (A^T * W * A)^{-1} * A^T * W * B$$




$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & -1 & 0 & 0 \\ 0 & 1 & -1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L_0 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \\ -10 \\ -5 \\ -1 \end{bmatrix}$$

# 1D Linear SLAM – case 4

Case 4 – Inconsistent Measurements with Confidence Matrix:

- Linear Least Squares allows us to include a weighting of each linear constraint.
- We can include weights in the computation
- We weight each constraint by a diagonal matrix where the weights are  $1/\text{variance}$  for each constraint.
- Let's say  $X_2$  **variance is 0.2**

$$X = (A^T * W * A)^{-1} * A^T * W * B$$



# 1D Linear SLAM – case 4


Case 4 – Inconsistent Measurements with Confidence Matrix:

- Linear Least Squares allows us to include a weighting of each linear constraint.
- We can include weights in the computation
- We weight each constraint by a diagonal matrix where the weights are  $1/\text{variance}$  for each constraint.
- Let's say  $X_2$  **variance is 0.2**

$$X = (A^T * W * A)^{-1} * A^T * W * B$$

$$x = \begin{bmatrix} -3 & 2.18 & 5.71 & 6.82 \end{bmatrix}$$

Why did the estimation just become worse??



# What about non-Linear Least Squares?

- Large number of geometric problems in computer vision are non-linear least-squares problems.

$$\mathbf{x} = \mathbf{h}(\theta)$$

where  $\mathbf{h} : \mathbf{R}^n \rightarrow \mathbf{R}^m$ .

- $\mathbf{x}$  is the measurement vector,  $\theta$  is the parameter vector.
- Write  $\mathbf{f}(\theta) = \mathbf{h}(\theta) - \mathbf{x}$ .
- We desire to minimize

$$\|\mathbf{f}(\theta)\|^2$$

over all choices of parameter  $\theta$ .

# Gauss Newton Solution

1. Start from an initial value  $\theta_0$ .
2. At step  $i$  assume a linear approximation for the function at  $\theta_i$

$$\mathbf{f}(\theta_i + \Delta) = \mathbf{f}(\theta_i) + \mathbf{f}_\theta \Delta \text{ where } \mathbf{f}_\theta = \partial \mathbf{f} / \partial \theta = \mathbf{J} .$$

3. Solve

$$\mathbf{f}(\theta_i + \Delta) = \mathbf{f}(\theta_i) + \mathbf{J} \Delta = 0$$

or

$$\mathbf{J} \Delta = -\mathbf{f}(\theta_i)$$


4. This is a linear least-squares problem (solve for  $\Delta$ ):

$$\mathbf{J}^\top \mathbf{J} \Delta = \mathbf{J}^\top \mathbf{f}(\theta_i)$$


5. Then set  $\theta_{i+1} = \theta_i + \Delta$ .

**Gauss-Newton update equation**

$$\mathbf{J}^\top \mathbf{J} \Delta = -\mathbf{J}^\top \mathbf{f}$$



1D Gauss-Newton (Newton)  
iteration.



1D Gauss-Newton (Newton)  
iteration (failure)




# Gradient Descent

Search direction is the direction of fastest descent of the function  $g$ .

## Gradient descent update equation

$$\lambda \Delta = -g_{\theta} = -\mathbf{J}^{\top} \mathbf{f}$$

Requires a 1D line search in  $\lambda$  to find the optimum direction.




# Levenberg-Marquadt


- Mixture of Gauss-Newton and Gradient descent.
  - Acts like Gauss-Newton when close to the minimum (quadratic region)
  - Gradient descent when improvement is difficult.
  - Depends on a parameter  $\lambda$  which
    1. Controls the mixture of Gauss-Newton and Gradient Descent
    2. Controls the step-length.
-

# What about non-Linear Least Squares?


- Lets See some examples 1:



Gauss-Newton




Gradient descent




Levenberg

# What about non-Linear Least Squares?


- Lets See some examples 2:



Gauss-Newton




Gradient descent




Levenberg

# What about non-Linear Least Squares?


- Lets See some examples 3:



Gauss-Newton



Gradient descent



Levenberg

- It is obvious that Levenberg Marquadt displays robustness

# Bundle Adjustment

- Bundle Adjustment is the employment of nonlinear optimization in the problem of the minimization of the re-projection error, by finding the optimal Poses (extrinsics) of the cameras and the locations of the 3D points.


$$\arg \min_{\mathbf{w}, \boldsymbol{\theta}} \sum_{f=1}^F \sum_{n=1}^N ||\mathbf{x}_n^f - \pi(\mathbf{w}_n; \boldsymbol{\theta}^f)||_2^2$$

$\mathbf{x}$  2D projection  $\mathbf{w}$  3D point

$\boldsymbol{\theta}$  extrinsics  $N$  no. of points

$\pi$  projection function

$F$  no. of frames



# Bundle Adjustment – Linearization

$$\pi(\mathbf{w}_n + \Delta \mathbf{w}_n; \boldsymbol{\theta}_f \circ \Delta \boldsymbol{\theta}_f) \approx \pi(\mathbf{w}_n; \boldsymbol{\theta}_f) + \mathbf{J}_n^f \begin{bmatrix} \Delta \boldsymbol{\theta}_f \\ \Delta \mathbf{w}_n \end{bmatrix}$$



$$\arg \min_{\Delta \boldsymbol{\theta}, \Delta \mathbf{w}} \sum_{f=1}^F \sum_{n=1}^N \rho_n^f \left\| \mathbf{x}_n^f - \pi(\mathbf{w}_n; \boldsymbol{\theta}_f) - \mathbf{J}_n^f \begin{bmatrix} \Delta \boldsymbol{\theta}_f \\ \Delta \mathbf{w}_n \end{bmatrix} \right\|_2^2$$

$\mathbf{x}$  2D projection  $\mathbf{w} \leftarrow$  3D point

$\boldsymbol{\theta}$  extrinsics  $N$  no. of points

$\pi$  projection function

$F$  no. of frames  $\rho \rightarrow$  visibility  $\in [0, 1]$

# Bundle Adjustment – Linearization


- The Linearization of the minimization happens by calculating the Jacobian of the projection matrix
- Assuming the following projection matrix: 
$$\begin{bmatrix} wu \\ wv \\ w \end{bmatrix} = \begin{bmatrix} f_x & s_k & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} [R \quad T] \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$
- We first define the orientation as the rotation matrix associated with the axis angle  $w_x, w_y, w_z$  using the Rodrigues equation
- The Jacobian **FUNCTION** can be calculated as:

$$\mathbf{J} = \begin{bmatrix} \frac{\partial u}{\partial w_x} & \frac{\partial u}{\partial w_y} & \frac{\partial u}{\partial w_z} & \frac{\partial u}{\partial f} & \frac{\partial u}{\partial u_0} & \frac{\partial u}{\partial v_0} & \frac{\partial u}{\partial X} & \frac{\partial u}{\partial Y} & \frac{\partial u}{\partial Z} \\ \frac{\partial v}{\partial w_x} & \frac{\partial v}{\partial w_y} & \frac{\partial v}{\partial w_z} & \frac{\partial v}{\partial f} & \frac{\partial v}{\partial u_0} & \frac{\partial v}{\partial v_0} & \frac{\partial v}{\partial X} & \frac{\partial v}{\partial Y} & \frac{\partial v}{\partial Z} \end{bmatrix}$$



# Bundle Adjustment – Comments

- Bundle adjustment (and graph optimization) is the backbone of all SLAM algorithms
- Keep in mind that:
  - We need to provide the Jacobian of the projection
  - We usually provide a covariance matrix (See the linear case for uncertainty)
  - It is solved using Levenberg-Marquadt
  - There are a lot of computational issues which are overcome exploiting the sparsity of the function  $AX=b$  (see least squares)  
Look at the following A and b Matrices  
This solution is called **Sparse Bundle Adjustment!**



# Bundle Adjustment – Covariance


- The uncertainty of the camera pose  $C_k$  is a combination of the uncertainty at  $C_{k-1}$  (black-solid ellipse) and the uncertainty of the transformation  $T_k$  (gray dashed ellipse)

- $C_k = f(C_{k-1}, T_k)$

- The combined covariance  $\Sigma_k$  is

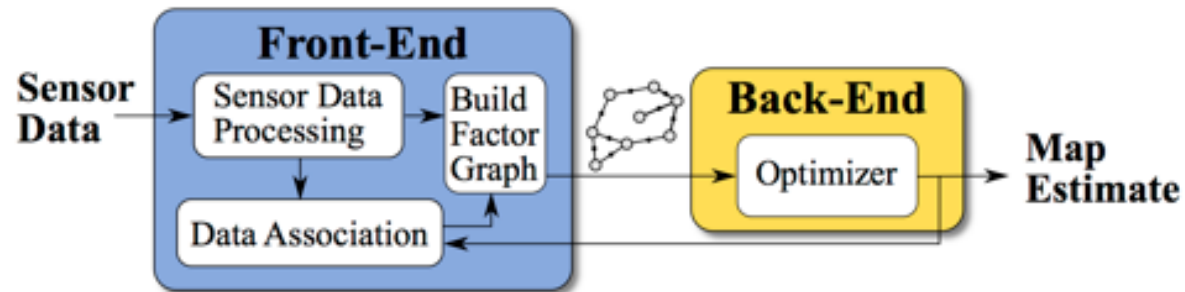
$$\begin{aligned}\Sigma_k &= J \begin{bmatrix} \Sigma_{k-1} & 0 \\ 0 & \Sigma_{k,k-1} \end{bmatrix} J^\top \\ &= J_{\vec{C}_{k-1}} \Sigma_{k-1} J_{\vec{C}_{k-1}}^\top + J_{\vec{T}_{k,k-1}} \Sigma_{k,k-1} J_{\vec{T}_{k,k-1}}^\top\end{aligned}$$

- The camera-pose uncertainty is always increasing when concatenating transformations. Thus, it is important to keep the uncertainties of the individual transformations small



# Recent Visual Slam Solutions - Intro

- That was too much info, let's see now some recent solutions to the Slam Problem:
- Most recent visual SLAM methods are split in two parts:
- The **Frontend**: where the raw data are converted into pose graphs and Loop constraints and the **Backend** where, given a graph with constraints, the new pose of the robot is calculated as well as the surrounding map points.



# Recent Visual Slam Solutions – Common Architecture

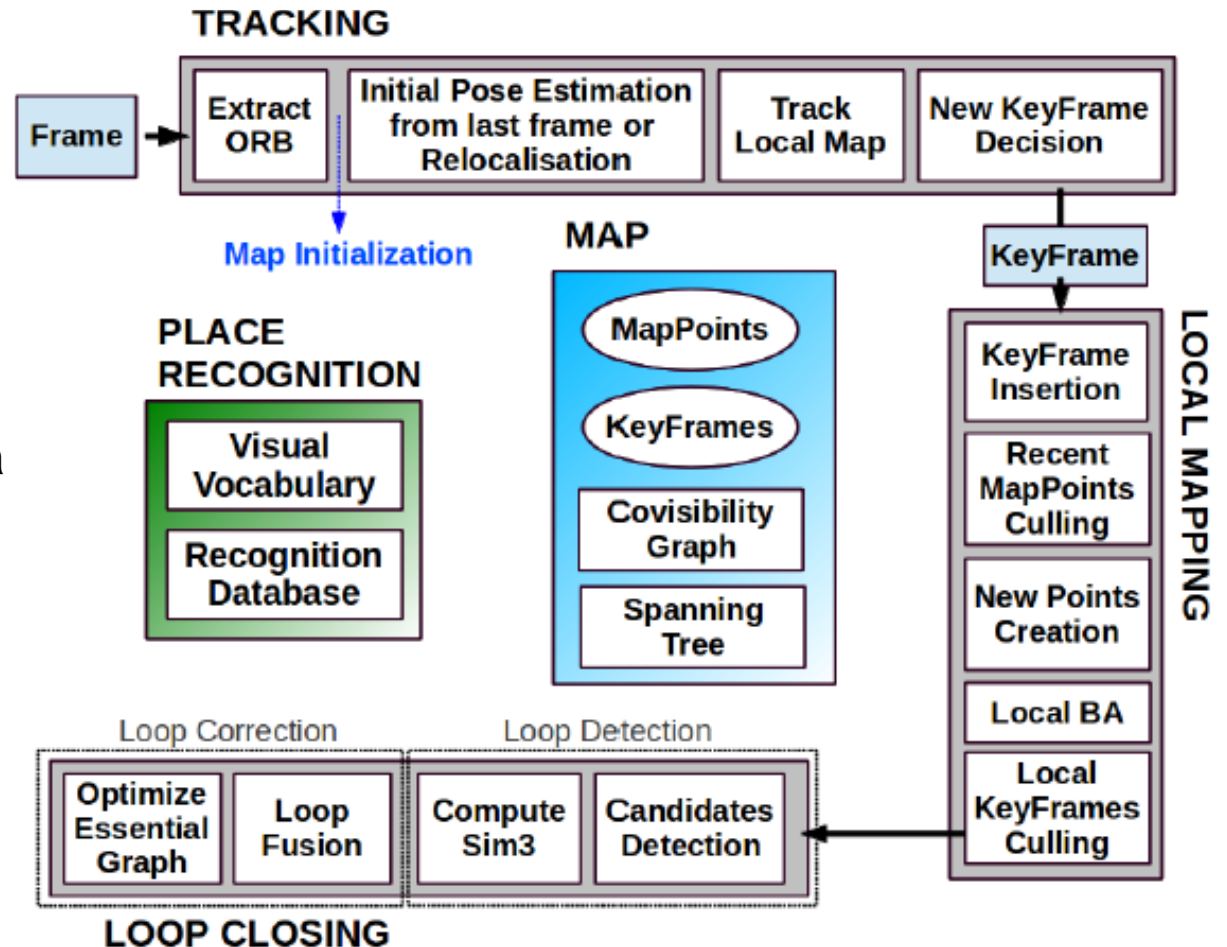
- Front End
  - Data Association
    - Frame to Frame
    - Multi-frame
    - Loop Closure Detection
  - Geometric Initialization
    - Pose Estimation
    - Landmark Triangulation
  - System Formation
    - Observation Matrix
    - Covariance Matrix
    - Graph Generation and Update
- Back End
  - Filter-Based State Estimation
    - Extended Kalman Filter
    - Particle Filters
  - Least squares optimization
    - Bundle Adjustment
    - Graph Optimization
      - Key Frame

# Recent Visual Slam Solutions – recent advances

Towards Realtime operation?:

- The computational cost of bundle adjustment has lead to the idea of **keyframing**:  
i.e.: identifying and describing some of the frames to be used for graph optimization.
- Bags of words for robust loop closure.
  - What can you tell me about that?
- Co-visibility Graph

- The ORBSLAM algorithm is one of the most well performing opensource implementations of visual slam.
- Three parallel threads:
  - tracking,
    - localizing the camera with every frame and deciding when to insert a new keyframe
  - local mapping
    - Processes new keyframes and performs local BA to achieve an optimal
- reconstruction in the surroundings of the camera
  - loop closing
    - The loop closing searches for loops with every new keyframe
    - Essential Graph





# ORB SLAM on Kitti

## ORB-SLAM

Raúl Mur-Artal, J. M. M. Montiel and Juan D. Tardós

{raulmur, josemari, tardos} @unizar.es



Instituto Universitario de Investigación  
en Ingeniería de Aragón  
**Universidad** Zaragoza



**Universidad**  
Zaragoza

# Sum up

- Sum up Localization from last time
- Some terminology
- Pose-Landmark Graph Slam
- Example of Linear 1D SLAM
- Non-Linear Optimization approaches
- Bundle Adjustment
- Visual Slam System architecture
- ORBSLAM



Perception for Autonomous Systems 31392:

# Visual SLAM

## *Simultaneous Localization & Mapping*

Lecturer: Evangelos Boukas—PhD